

Floodlight Functionalities By Examples


Carlo Vallati

Assistant Professor@ University of Pisa

c.vallati@iet.unipi.it

Hub

- Retransmit the packet on all the ports of the switch
- Create a module that transforms the SDN switches into hubs that retransmit all the packets in all the ports



- All the packets are processed at the controller
- No new rules are installed, every time a new Packet-Out is created to send each packet to the controller

```
// Cast to Packet-In
OFPacketIn pi = (OFPacketIn) msg;

// Create the Packet-Out and set basic fields
OFPacketOut.Builder pob = sw.getOFFactory().buildPacketOut();
pob.setBufferId(pi.getBufferId());

// Create action -> flood the packet on all the ports
OFActionOutput.Builder actionBuilder = sw.getOFFactory().actions().buildOutput();
actionBuilder.setPort(OFPort.FLOOD);

// Assign the action
pob.setActions(Collections.singletonList((OFAction) actionBuilder.build()));
```

- Packet waiting for being transmitted can be sent to the Controller encapsulated in the Packet-In message or can be buffered on the switch

```
// Packet might be buffered in the switch or encapsulated in Packet-In  
// If the packet is encapsulated in Packet-In sent it back  
if (pi.getBufferId() == OFBufferId.NO_BUFFER) {  
    // Packet-In buffer-id is none, the packet is encapsulated -> send it back  
    byte[] packetData = pi.getData();  
    pob.setData(packetData);  
}  
// Send the Packet-Out  
sw.write(pob.build());  
  
// Interrupt the chain  
return Command.STOP;
```

Generate Flow-Mods

- Instead of processing all the packets at the controller a Flow-Mod message can be generated to install new rules in the switches to flood all the packets in all the other ports

```
// Create a new rule to be added
OFFlowAdd.Builder fmb = sw.getOFFactory().buildFlowAdd();
fmb.setBufferId(pi.getBufferId()) // Link the new rule to the received OF PKT IN
    .setHardTimeout(20) // Set hard timeout
    .setIdleTimeout(10) // Set soft timeout
    .setPriority(32768) // Set priority
    .setXid(pi.getXid());


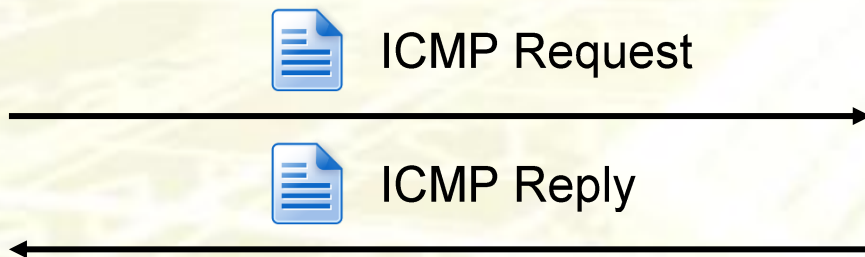
// Create a new action to be executed
OFActionOutput.Builder actionBuilder = sw.getOFFactory().actions().buildOutput();
actionBuilder.setPort(OFPort.FLOOD); // Set as action the flood of the packet
fmb.setActions(Collections.singletonList((OFAction) actionBuilder.build()));

// Send the Packet-Mod
sw.write(fmb.build());

// Interrupt the chain
return Command.STOP;
```

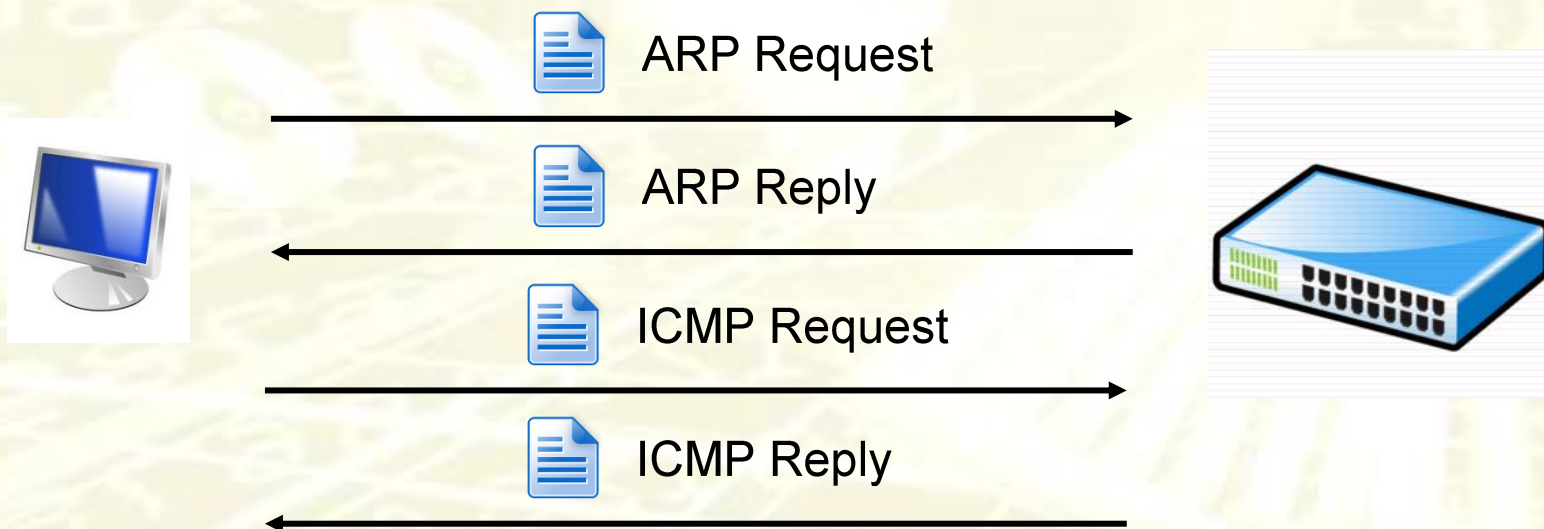

Virtual address

- Create a module that mimic the existence of Virtual hosts processing some traffic directed to a *virtual IPv4 address*.
- Specifically, the controller has to process ICMP ping requests to a specific IP address (that does not physically exist) generating ICMP replies



Virtual address

- No new rules are installed, every time a new Packet-Out is created to send each packet to the controller
- To this aim, the controller has also to handle ARP requests. Specifically the controller has to reply to ARP requests aimed at the virtual IPv4 address



Mininet



- If the virtual address is an address outside the local network of hosts emulated in mininet an extra configuration on them is required (they do not have a gateway configured)
- Specifically, each host has to be configured to send out traffic for external networks on its NIC card as it is
 - `h1 route add -net 0.0.0.0/32 dev h1-eth0`

Virtual address

- Inside the receiver function the controller has to process ARP requests and ICMP echo requests for the virtual IPv4:

```
// Cast packet
OFPacketIn pi = (OFPacketIn) msg;
// Dissect Packet included in Packet-In

IPacket pkt = eth.getPayload();

if (eth.isBroadcast() || eth.isMulticast()) {
    if (pkt instanceof ARP) {
        handleARPRequest(); // Handle ARP requests for the virtual IP
    }
}

// We only care about packets which are sent to the virtual IP address
IPv4 ip_pkt = (IPv4) pkt;

if(ip_pkt.getDestinationAddress().compareTo(VIRTUAL_IP) != 0){

    // Allow the next module to also process this OpenFlow message
    return Command.CONTINUE;
}

handleIPPacket(); // Handle IP packets towards the Virtual IP

// Do not continue processing this OpenFlow message
return Command.STOP;
```

Virtual address - handleARPRequest

- In case the ARP request is directed to the virtual IP an ARP reply with a virtual MAC address has to be generated

```
// Double check that the payload is ARP
Ethernet eth = IFloodlightProviderService.bcStore.get(cntx,
    IFloodlightProviderService.CONTEXT_PI_PAYLOAD);
if (! (eth.getPayload() instanceof ARP))
    return;
// Cast the ARP request
ARP arpRequest = (ARP) eth.getPayload();
// Generate ARP reply
IPacket arpReply = new Ethernet()
    .setSourceMACAddress(VIRTUAL_MAC)
    .setDestinationMACAddress(eth.getSourceMACAddress())
    .setEtherType(EthType.ARP)
    .setPriorityCode(eth.getPriorityCode())
    .setPayload(
        new ARP()
            .setHardwareType(ARP.HW_TYPE_ETHERNET)
            .setProtocolType(ARP.PROTO_TYPE_IP)
            .setHardwareAddressLength((byte) 6)
            .setProtocolAddressLength((byte) 4)
            .setOpCode(ARP.OP_REPLY)
            .setSenderHardwareAddress(VIRTUAL_MAC) // Set my MAC address
            .setSenderProtocolAddress(VIRTUAL_IP) // Set my IP address
            .setTargetHardwareAddress(arpRequest.getSenderHardwareAddress())
            .setTargetProtocolAddress(arpRequest.getSenderProtocolAddress())
    );
```

Virtual address - handleARPRequest

- Create Packet-Out
- Create list of actions
- Set the ARP reply as payload of Packet-Out
- Send it out

// Initialize a packet out

```
OFPacketOut.Builder pob = sw.getOFFactory().buildPacketOut();  
pob.setBufferId(OFBufferId.NO_BUFFER);  
pob.setInPort(OFPort.ANY);
```

// Set the output action

```
OFActionOutput.Builder actionBuilder = sw.getOFFactory().actions().buildOutput();  
OFPort inPort = pi.getMatch().get(MatchField.IN_PORT);  
actionBuilder.setPort(inPort);  
pob.setActions(Collections.singletonList((OFAction) actionBuilder.build()));
```

// Set the ARP reply as packet data

```
byte[] packetData = arpReply.serialize();  
pob.setData(packetData);
```

// Send packet

```
sw.write(pob.build());
```

Virtual address - handleIPPacket

- Check Packet-In
- Create the ICMP reply


```
// Cast the IP packet
IPv4 ipv4 = (IPv4) eth.getPayload();

// Check that the IP is actually an ICMP request
if (! (ipv4.getPayload() instanceof ICMP))
    return;

// Cast to ICMP packet
ICMP icmpRequest = (ICMP) ipv4.getPayload();
// Generate ICMP reply
IPacket icmpReply = new Ethernet()
    .setSourceMACAddress(VIRTUAL_MAC)
    .setDestinationMACAddress(eth.getSourceMACAddress())
    .setEtherType(Ethernet.TYPE_IPv4)
    .setPriorityCode(eth.getPriorityCode())
    .setPayload(
        new IPv4()
            .setProtocol(IpProtocol.ICMP)
            .setDestinationAddress(ipv4.getSourceAddress())
            .setSourceAddress(VIRTUAL_IP)
            .setTtl((byte) 64)
            .setProtocol((byte) 4)
            // Set the same payload included in the request
            .setPayload(
                new ICMP()
                    .setIcmpType(ICMP.ECHO_REPLY)
                    .setIcmpCode(icmpRequest.getIcmpCode())
                    .setPayload(icmpRequest.getPayload())
            )
    );
```


Load Balancer

- Create a module that acts as a load balancer for a pool of servers
- A virtual IP address (e.g. 8.8.8.8) is set as the public IP address of a group of servers offering the same service (e.g. a web server with the same page)
- The module has to control the switch to dispatch the requests directed to the virtual IP to one of the IP of a real server from the pool.
- The real server from the pool selected to provide the service must change every 20 seconds in order to balance the load



Load Balancer

- *The switch has to handle ARP requests.*
- *This operation has to be performed transparently*, i.e. changing the destination IP and MAC addresses on the requests (with the addresses of the real server selected) and the source IP and MAC addresses on the responses (with the addresses of the virtual IP and MAC)
- *New rules have to be installed on the switch*



Load Balancer



// Create a flow table modification message to add a rule

```
OFFlowMod.Builder fmb = sw.getOFFactory().buildFlowAdd();
```

```
fmb.setIdleTimeout(IDLE_TIMEOUT);
```

```
fmb.setHardTimeout(HARD_TIMEOUT); // Set as hard timeout the period to change the current server
```

```
fmb.setBufferId(OFFBufferId.NO_BUFFER);
```

```
fmb.setOutPort(OFFPort.CONTROLLER);
```

```
fmb.setBufferId(pi.getBufferId());
```

```
fmb.setCookie(U64.of(0));
```

```
fmb.setPriority(FlowModUtils.PRIORITY_MAX);
```

// Create the match for the new rule (incoming IPv4 traffic directed to the load balancer)

```
Match.Builder mb = sw.getOFFactory().buildMatch();
```

```
mb.setExact(MatchField.ETH_TYPE, EthType.IPv4)
```

```
.setExact(MatchField.IPV4_DST, LOAD_BALANCER_IP)
```

```
.setExact(MatchField.ETH_DST, LOAD_BALANCER_MAC);
```

Load Balancer



```
// Create the list of actions associated with a match
OFActions actions = sw.getOFFactory().actions();
ArrayList<OFAction> actionList = new ArrayList<OFAction>();
OFOxms oxms = sw.getOFFactory().oxms();
// Set as new MAC destination the MAC address of the current server
OFActionSetField setDlDst = actions.buildSetField()
    .setField(
        oxms.buildEthDst()
        .setValue(MacAddress.of(SERVERS_MAC[last_server]))
        .build()
    )
    .build();
actionList.add(setDlDst);
// Set as new IP destination the IP address of the current server
OFActionSetField setNwDst = actions.buildSetField()
    .setField(
        oxms.buildIpv4Dst()
        .setValue(IPv4Address.of(SERVERS_IP[last_server]))
        .build()
    ).build();
actionList.add(setNwDst);
// Set as output port the port of the current server
OFActionOutput output = actions.buildOutput()
    .setMaxLen(0xFFffFFff)
    .setPort(OFFPort.of(SERVERS_PORT[last_server]))
    .build();
actionList.add(output);

// Send out the mod message
fmb.setActions(actionList);
fmb.setMatch(mb.build());

sw.write(fmb.build());
```

References

- *Google, SDN experience:*
<http://www.opennetsummit.org/archives/apr12/hoelzle-tue-openflow.pdf>
- *Mininet, walkthrough:*
<http://mininet.org/walkthrough/>
- *Floodlight documentation:*
<http://www.projectfloodlight.org/documentation/>
- *OpenFlow tutorial:*
http://archive.openflow.org/wk/index.php/OpenFlow_Tutorial